

Machine Learning

The Maximum Entropy classifier

Maximum Entropy Model

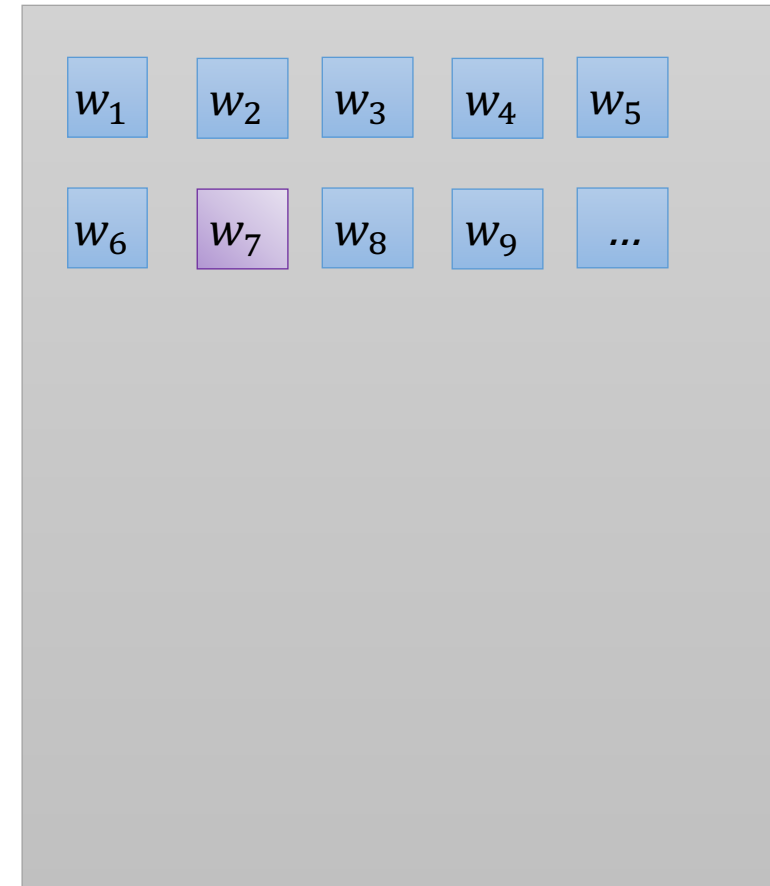
- We are now deriving one of most widely used classifier (still today!)
- The intuition is based on the „*Principle of Maximum Entropy*“

Principle of Maximum Entropy (Jaynes, 1957)

If incomplete information about a probability distribution is available, the only unbiased assumption that can be made is a distribution which is as uniform as possible, given the available information.

Principle of Maximum Entropy

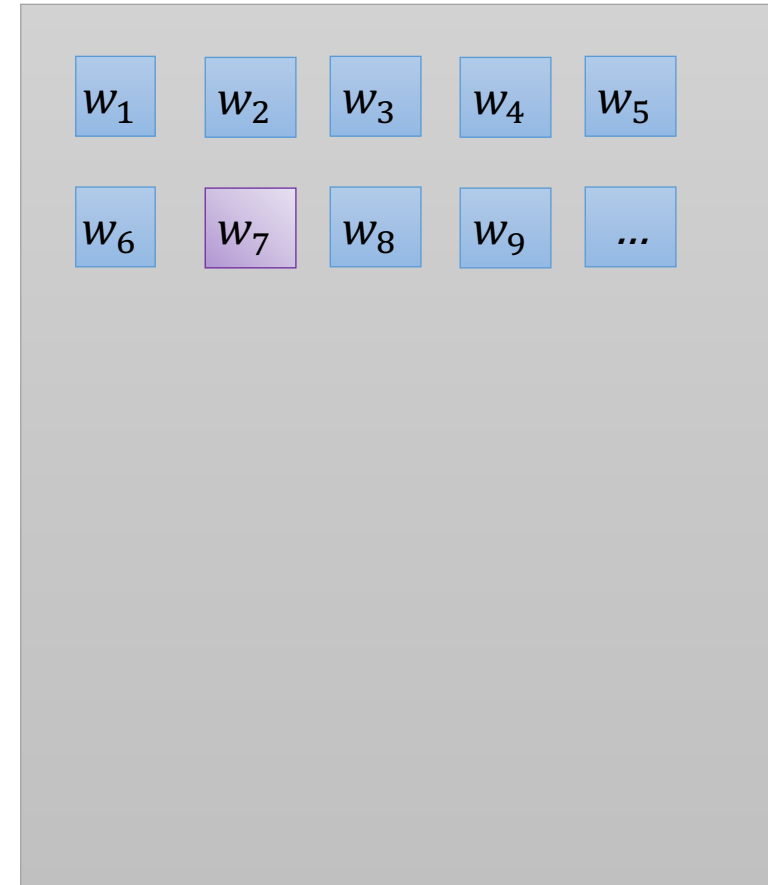
- What does this even mean?
- Assume, that we are given text and want to predict POS-tags for every word
- And we know (from arbitrary sources, e.g. Dbpedia, labelled data), that „Markus“ is 4 times as likely to be a noun than a verb
- Since we have no idea about the rest, we would assume they are uniformly distributed(all are equally as probable)



Principle of Maximum Entropy


- Introduction of features
- We can formulate the piece of knowledge, that „the word = Markus is 4 times more likely to be a noun than a verb“ using **feature functions**
- Let us introduce the feature:
 - Currentword = „Markus“

We will address this soon!



Principle of Maximum Entropy

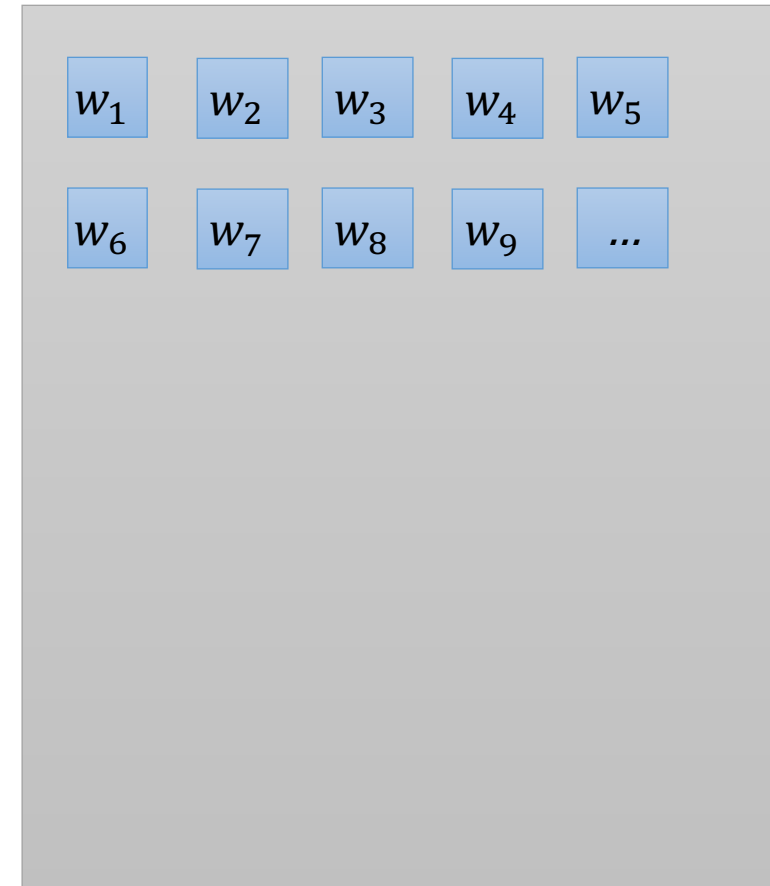
- In the same manner, we can introduce arbitrary features:
 - Previous-Word = „Markus“
 - Next-Word = „Markus“
 - Previous-Word = „and“
 - Next-Word = „drives“
 - ...
- We could in fact use whatever comes to our mind



Features in Maximum Entropy

- The introduction of features was a genius idea! We can now describe all of our words using a set of features
- Let us do this while making use of three simple features:

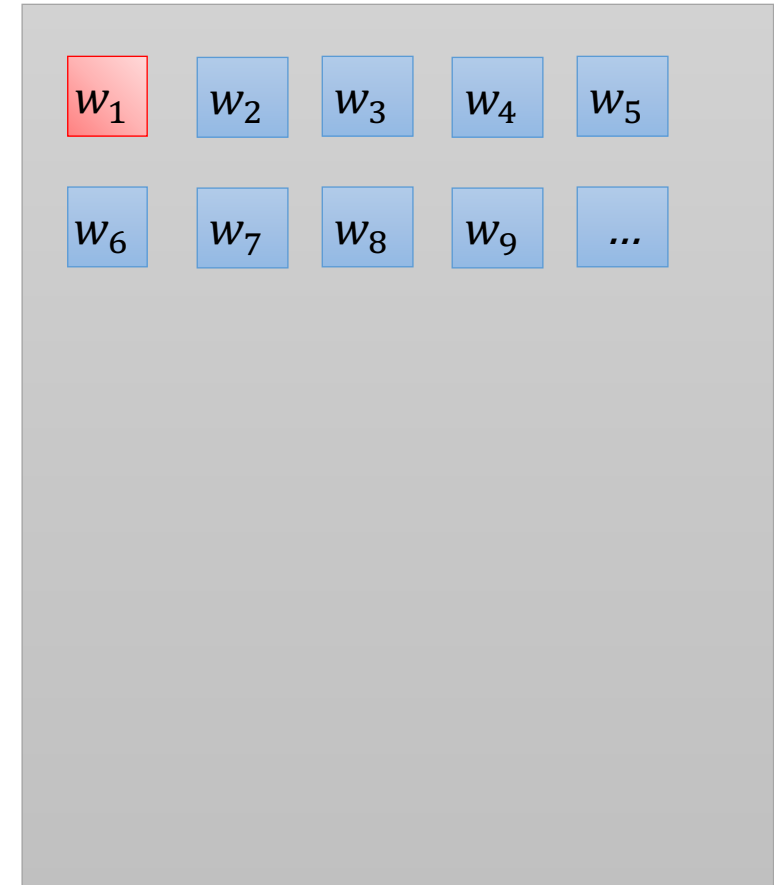
Feature-Name	Feature Value
Current Word	
Previous Word	
Next Word	



Features in Maximum Entropy

- We can now describe our words as follows
- For w_1

Feature-Name	Feature Value
Current Word	w_1
Previous Word	start
Next Word	w_2

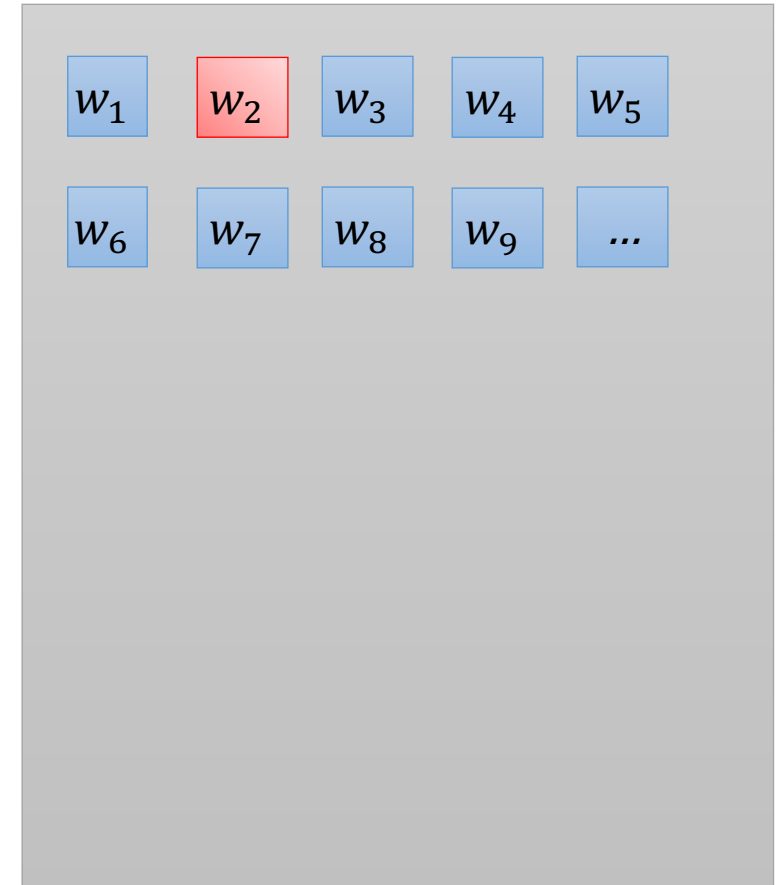


Features in Maximum Entropy

- We can now describe our words as follows
- For w_2

Feature-Name	Feature Value
Current Word	w_2
Previous Word	w_1
Next Word	w_3

... and so on for the other words



Features in Maximum Entropy

- After this point it is no longer important, that we started with a text classification process, since all we have left are tables of features for every instance (a word)



Feature-Name [ID]	Feature Value
Current Word	w_1
Previous Word	start
Next Word	w_2

Feature-Name [ID]	Feature Value
Current Word	w_2
Previous Word	w_1
Next Word	w_3

... more tables

Features in Maximum Entropy

- We did not deal with one aspect:
- „the word = Markus is 4 times more likely to be a noun than a verb“
- Let us aggregate all our tables into one huge table

Feature-Name [ID]	Feature Value
Current Word	w_1
Previous Word	start
Next Word	w_2

Feature-Name [ID]	Feature Value
Current Word	w_2
Previous Word	w_1
Next Word	w_3



Feature-Name = Value
Current Word = w_1
Current Word = w_2
Previous Word = start
Previous Word = w_1
Next Word = w_2
Next Word = w_3
... many many more

Features in Maximum Entropy

- We did not deal with one aspect:
- „the word = Markus is 4 times more likely to be a noun than a verb“
- Let us aggregate all our tables into one huge table


Feature-Name = Value	Active
Current Word = w_1	1
Current Word = w_2	0
Previous Word = start	1
Previous Word = w_1	0
Next Word = w_2	1
Next Word = w_3	0



Features in Maximum Entropy

- We did not deal with one aspect:
- „the word = Markus is 4 times more likely to be a noun than a verb“
- Let us aggregate all our tables into one huge table

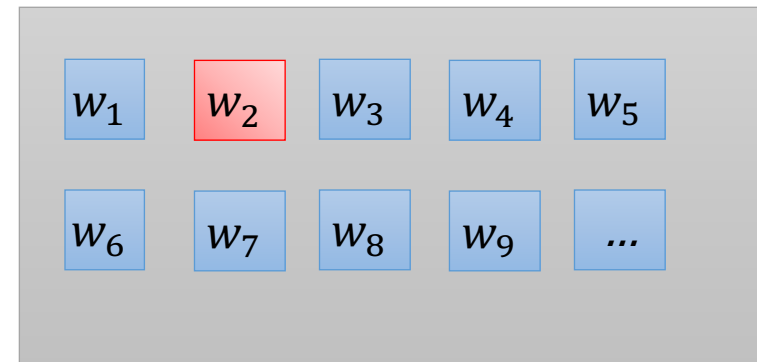
Feature-Name = Value	Active
Current Word = w_1	0
Current Word = w_2	1
Previous Word = start	0
Previous Word = w_1	1
Next Word = w_2	0
Next Word = w_3	1



Features in Maximum Entropy

- We can now describe our entire data set using a common vector of these features, but how does this help to express „more likely“

Feature-Name = Value	Active
Current Word = w_1	0
Current Word = w_2	1
Previous Word = start	0
Previous Word = w_1	1
Next Word = w_2	0
Next Word = w_3	1

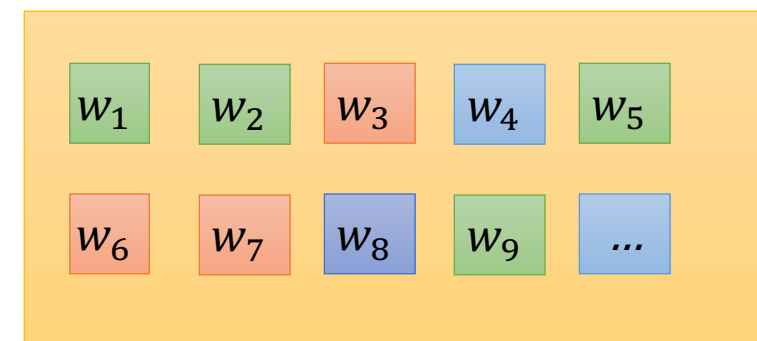





Features in Maximum Entropy

- Introduce labelled data
- We can now also encode the labels into our features

Feature-Name = Value	Active
Current Word = w_1	0
Current Word = w_2	1
Previous Word = start	0
Previous Word = w_1	1
Next Word = w_2	0
Next Word = w_3	1


Legend: Colors of labels



Features in Maximum Entropy

- Introduce labelled data
- We can now also encode the labels into our features

Feature-Name = Value	Active
Current Word = w_1	0
Current Word = w_2	1
Previous Word = start	0
Previous Word = w_1	1
Next Word = w_2	0
Next Word = w_3	1



Feature-Name = Value	Current-Label	Active
Current Word = w_1	ADJ	
Current Word = w_1	N	
Current Word = w_1	V	
Current Word = w_2	ADJ	
Current Word = w_2	N	
Current Word = w_2	V	
Previous Word = w_1	ADJ	
Previous Word = w_1	N	
Previous Word = w_1	V	
...		

Features in Maximum Entropy

- Let us redescribe our words
- For w_1

Feature-Name = Value	Current- Label	Active
Current Word = w_1	ADJ	1
Current Word = w_1	N	0
Current Word = w_1	V	0
Current Word = w_2	ADJ	0
Current Word = w_2	N	0
Current Word = w_2	V	0
Previous Word = w_1	ADJ	0
Previous Word = w_1	N	0
Previous Word = w_1	V	0
...		



Features in Maximum Entropy

- Let us redescribe our words
- For w_2

Feature-Name = Value	Current- Label	Active
Current Word = w_1	ADJ	0
Current Word = w_1	N	0
Current Word = w_1	V	0
Current Word = w_2	ADJ	1
Current Word = w_2	N	0
Current Word = w_2	V	0
Previous Word = w_1	ADJ	1
Previous Word = w_1	N	0
Previous Word = w_1	V	0
...		



Features in Maximum Entropy

- Let us now sum all these tables of all words

Feature-Name = Value	Current- Label	Active
Current Word = w_1	ADJ	1
Current Word = w_1	N	80
Current Word = w_1	V	20
Current Word = w_2	ADJ	43
Current Word = w_2	N	12
Current Word = w_2	V	1337
Previous Word = w_1	ADJ	42
Previous Word = w_1	N	11
Previous Word = w_1	V	0
...		

Features in Maximum Entropy

- Assume $w_1 = \text{Markus}$
- We would expect to have it labelled as a noun N
- Our belief is roughly

$$p(N|w = \text{Markus}) \sim \frac{80}{1 + 80 + 20} \sim 80\%$$

$$p(V|w = \text{Markus}) \sim \frac{20}{1 + 80 + 20} \sim 20\%$$

We did it! We now have an intuition how we can express something to be more likely than something else, just using features

Feature-Name = Value	Current- Label	Active
Current Word = w_1	ADJ	1
Current Word = w_1	N	80
Current Word = w_1	V	20
Current Word = w_2	ADJ	43
Current Word = w_2	N	12
Current Word = w_2	V	1337
Previous Word = w_1	ADJ	42
Previous Word = w_1	N	11
Previous Word = w_1	V	0
...		

Features in Maximum Entropy

Principle of Maximum Entropy (Jaynes, 1957)

If **incomplete information about a probability distribution** is available, the only unbiased assumption that can be made is a distribution which is as uniform as possible, **given the available information**.

Even though we counted 0, this might be due to our limited amount of training data, so we assume that we do not know anything about this!

Feature-Name = Value	Current- Label	Active
Current Word = w_1	ADJ	1
Current Word = w_1	N	80
Current Word = w_1	V	20
Current Word = w_2	ADJ	43
Current Word = w_2	N	12
Current Word = w_2	V	1337
Previous Word = w_1	ADJ	42
Previous Word = w_1	N	11
Previous Word = w_1	V	0
...		

Features in Maximum Entropy

- Now we want to learn a classifier that is exactly able to reproduce this table
 - We will call such a classifier „consistent“ to our observations!

Feature-Name = Value	Current- Label	Active
Current Word = w_1	ADJ	1
Current Word = w_1	N	80
Current Word = w_1	V	20
Current Word = w_2	ADJ	43
Current Word = w_2	N	12
Current Word = w_2	V	1337
Previous Word = w_1	ADJ	42
Previous Word = w_1	N	11
Previous Word = w_1	V	0
...		

Features in Maximum Entropy

- Let us merge some columns again
- Each row now corresponds to a distinct feature $f(x, y)$ which is calculated from the input x and the label y

Feature-Name = Value \wedge CurrentLabel	Active
Current Word = $w_1 \wedge \text{ADJ}$	1
Current Word = $w_1 \wedge \text{N}$	80
Current Word = $w_1 \wedge \text{V}$	20
Current Word = $w_2 \wedge \text{ADJ}$	43
Current Word = $w_2 \wedge \text{N}$	12
Current Word = $w_2 \wedge \text{V}$	1337
Previous Word = $w_1 \wedge \text{ADJ}$	42
Previous Word = $w_1 \wedge \text{N}$	11
Previous Word = $w_1 \wedge \text{V}$	0
...	

Features in Maximum Entropy

- We can now represent all data in a **single vector**! And the **semantics of each index** can be stored somewhere else

Feature-Name = Value \wedge CurrentLabel
Current Word = $w_1 \wedge$ ADJ
Current Word = $w_1 \wedge$ N
Current Word = $w_1 \wedge$ V
Current Word = $w_2 \wedge$ ADJ
Current Word = $w_2 \wedge$ N
Current Word = $w_2 \wedge$ V
Previous Word = $w_1 \wedge$ ADJ
Previous Word = $w_1 \wedge$ N
Previous Word = $w_1 \wedge$ V
...

Active
1
80
20
43
12
1337
42
11
0

Features in Maximum Entropy

- A training example is also just such a vector

Feature-Name = Value	Active
Current Word = $w_1 \wedge \text{ADJ}$	1
Current Word = $w_1 \wedge \text{N}$	0
Current Word = $w_1 \wedge \text{V}$	0
Current Word = $w_2 \wedge \text{ADJ}$	0
Current Word = $w_2 \wedge \text{N}$	0
Current Word = $w_2 \wedge \text{V}$	0
Previous Word = $w_1 \wedge \text{ADJ}$	0
Previous Word = $w_1 \wedge \text{N}$	0
Previous Word = $w_1 \wedge \text{V}$	0
...	

Features in Maximum Entropy

- Let us inspect the table a little more

Feature-Name = Value	Active
Current Word = $w_1 \wedge \text{ADJ}$	1
Current Word = $w_1 \wedge \text{N}$	0
Current Word = $w_1 \wedge \text{V}$	0
Current Word = $w_2 \wedge \text{ADJ}$	0
Current Word = $w_2 \wedge \text{N}$	0
Current Word = $w_2 \wedge \text{V}$	0
Previous Word = $w_1 \wedge \text{ADJ}$	0
Previous Word = $w_1 \wedge \text{N}$	0
Previous Word = $w_1 \wedge \text{V}$	0
...	

Only one of each
group can be active

Features in Maximum Entropy

- We can also write a probability distribution to each group

Feature-Name = Value	Active
Current Word = $w_1 \wedge \text{ADJ}$	1
Current Word = $w_1 \wedge \text{N}$	0
Current Word = $w_1 \wedge \text{V}$	0
Current Word = $w_2 \wedge \text{ADJ}$	0
Current Word = $w_2 \wedge \text{N}$	0
Current Word = $w_2 \wedge \text{V}$	0
Previous Word = $w_1 \wedge \text{ADJ}$	0
Previous Word = $w_1 \wedge \text{N}$	0
Previous Word = $w_1 \wedge \text{V}$	0
...	

100%
0%
0%

We observe this feature with ADJ, since it is labelled that way!

→ If a classifier is not certain about this labeling, its distribution would be somewhat different

Features in Maximum Entropy

- We can also write a probability distribution to each group

Feature-Name = Value	Active	Classifier
Current Word = $w_1 \wedge \text{ADJ}$	1	0.8
Current Word = $w_1 \wedge \text{N}$	0	0.1
Current Word = $w_1 \wedge \text{V}$	0	0.1
Current Word = $w_2 \wedge \text{ADJ}$	0	
Current Word = $w_2 \wedge \text{N}$	0	
Current Word = $w_2 \wedge \text{V}$	0	
Previous Word = $w_1 \wedge \text{ADJ}$	0	
Previous Word = $w_1 \wedge \text{N}$	0	
Previous Word = $w_1 \wedge \text{V}$	0	
...		

Classifier believes that this example is labelled ADJ with 80% certainty; 10% Noun and 10% Verb

Features in Maximum Entropy

- So the classifier provides us with his „soft counts“ (in fact these are expectations)

Feature-Name = Value	Active	Classifier
Current Word = $w_1 \wedge \text{ADJ}$	1	0.8
Current Word = $w_1 \wedge \text{N}$	0	0.1
Current Word = $w_1 \wedge \text{V}$	0	0.1
Current Word = $w_2 \wedge \text{ADJ}$	0	
Current Word = $w_2 \wedge \text{N}$	0	
Current Word = $w_2 \wedge \text{V}$	0	
Previous Word = $w_1 \wedge \text{ADJ}$	0	
Previous Word = $w_1 \wedge \text{N}$	0	
Previous Word = $w_1 \wedge \text{V}$	0	
...		

Features in Maximum Entropy

- So we can now produce two vectors, one with observed counts and one with „predicted counts“

Feature-Name = Value \wedge CurrentLabel
Current Word = $w_1 \wedge$ ADJ
Current Word = $w_1 \wedge$ N
Current Word = $w_1 \wedge$ V
Current Word = $w_2 \wedge$ ADJ
Current Word = $w_2 \wedge$ N
Current Word = $w_2 \wedge$ V
Previous Word = $w_1 \wedge$ ADJ
Previous Word = $w_1 \wedge$ N
Previous Word = $w_1 \wedge$ V
...

Observed
1
80
20
43
12
1337
42
11
0
...

Predicted
0.8
56.4
13.4
22.9
9.4
1228.8
56.3
14.5
0.5
...

Features in Maximum Entropy

- Goal now is to train a classifier, which is capable of reproducing the observed counts
- In fact we will derive that the difference between these 2 vectors corresponds to the gradient of the „MaximumEntropy“ classifier
 - And the gradient is all we need to use Gradient Descent!

Observed	Predicted
1	0.8
80	56.4
20	13.4
43	22.9
12	9.4
1337	1228.8
42	56.3
11	14.5
0	0.5

Maximum Entropy Model

- A MaxEnt classifier predicts a probability for a label y , given some feature vector x : $p(y|x)$
- Starting with the (Conditional Entropy)

$$H(Y|X) = - \sum_{(x,y) \in D} p(y, x) \log p(y|x)$$

- Our goal is to find a model $p^*(y|x)$, which:
 1. Maximizes the Conditional Entropy $H(Y|X)$ (Maximum Entropy)
 2. Is consistent to the training data (we will come to that in a moment)

Maximum Entropy Model: Goal

$$p^*(y|x) = \max_{p(y|x) \in \mathcal{P}} H(Y|X)$$

→ With \mathcal{P} being the set of consistent models (still to be modelled)

- We call a model consistent to our data, if it is consistent with respect to its “features”
 - Introduce “features” or “feature functions” :

$$f_i(x, y) = \begin{cases} 1, & \text{if a boolean expression is true} \\ 0, & \text{otherwise} \end{cases}$$

- Example:

$$f_i(x, y) = \begin{cases} 1, & \text{if } \text{currentWord} = \text{running and } y = \text{verb} \\ 0, & \text{otherwise} \end{cases}$$

Maximum Entropy Model: Consistency

- But what is consistency with respect to a feature ?
- For every feature function f_i , we can count how often it appears in our training data, and we can receive an expected value $\hat{E}(f_i)$:

$$\hat{E}(f_i) = \frac{1}{N} \sum_{(x,y) \in D} f_i(x, y)$$

- A valid model is a model, that can reproduce this expected value

Maximum Entropy Model: Consistency

- We assumed, that all training instances x are independent (and therefore equally probable), in general we would write:

$$\hat{E}(f_i) = \sum_{(x,y) \in D} p(x,y) \cdot f_i(x,y)$$

- We will now derive the same expected value for the model distribution, we will denote it as $E(f_i)$

Maximum Entropy Model: Consistency

- Our model predicts $p(y|x)$ so we have to do some work!

$$E(f_i) = \sum_{(x,y) \in D} p(x, y) \cdot f_i(x, y)$$

$$E(f_i) = \sum_{(x,y) \in D} p(x) \cdot p(y|x) \cdot f_i(x, y) \quad \text{Bayes}$$

- But we will yet again assume our instances are independant, so $p(x)$ is the same for every instance, and we can average again

$$E(f_i) = \frac{1}{N} \sum_{(x,y) \in D} p(y|x) \cdot f_i(x, y)$$

Maximum Entropy Model

- We can now force (for every feature), that
$$\hat{E}(f_i) = E(f_i)$$
- Every model that satisfies this is valid according to our observed data!
- We furthermore want to have real probability distribution:

$$p(y|x) \geq 0, \text{ for all } x, y$$

And

$$\sum_{y \in Y} p(y|x) = 1, \text{ for all } x$$

Maximum Entropy Model

- The full story:
 - Find $p(y|x)$ so that:

$$p^*(y|x) = \max_{p(y|x) \in P} H(Y|X)$$

Subject to:

$$\hat{E}(f_i) - E(f_i) = 0, \quad \forall f_i$$

$$1 - \sum_{y \in Y} p(y|x) = 0, \quad \forall x$$

And

$$p(y|x) \geq 0, \quad \forall x, y$$

Maximum Entropy Model

- Solving the optimization problem
 - Optimization without conditions is easy
 - But there is a mathematical trick to get rid of them
 - ➔ Introduce Lagrangian parameter

$$p^*(y|x) = \max_{p(y|x) \in P} H(Y|X)$$

Subject to:

$$\hat{E}(f_i) - E(f_i) = 0, \quad \forall f_i$$

$$1 - \sum_{y \in Y} p(y|x) = 0, \quad \forall x$$

And


$$p(y|x) \geq 0, \quad \forall x, y$$

Lagrange function:

$$L(p, \vec{\lambda}) = \underbrace{H(y|x)}_{\text{Entropie}} + \underbrace{\sum_i \lambda_i (E(f_i) - \hat{E}(f_i))}_{\text{Valid models}} + \underbrace{\lambda_{m+1} (1 - \sum_{y \in Y} p(y_i|x))}_{\text{Prob. Distribution}}$$

➔ We can now maximize this function instead!

Getting to the solution



Maximum Entropy Model

- Solution: The general form of a maximum entropy model (a “family” of models)

$$p(y|x) = \frac{\exp \sum_{i=1} \lambda_i f_i(x, y)}{\sum_{y' \in Y} \exp(\sum_{i=1} \lambda_i f_i(x, y'))}$$

- With model parameters λ_i , that are optimized for every task!

Maximum Entropy Model

- What you should understand until now:
 - Beginning with the principle of Maximum Entropy, we found a family of functions (“Exponential Family”) that can be used to classify!

$$p(y|x) = \frac{\exp \sum_{i=1} \lambda_i f_i(x, y)}{\sum_{y' \in Y} \exp(\sum_{i=1} \lambda_i f_i(x, y'))}$$

- You do not know anything about how to get those λ_i !

Maximum Entropy Model

- Let us assume all λ_i are 0

$$p(y|x) = \frac{\exp \sum_{i=1} \lambda_i f_i(x, y)}{\sum_{y' \in Y} \exp(\sum_{i=1} \lambda_i f_i(x, y'))} = \frac{1}{\sum_{y' \in Y} 1} = \frac{1}{\# \text{ labels}}$$

- ➔ If we have no information, we will just predict a uniform distribution! (Principle of Maximum Entropy)
- ➔ The same would hold if no feature function would be active!

Maximum Entropy Model: Example

- Problem: Determine the POS-Tag for the word “**Grace**” in the sentence:

“at **Grace** Road”

- We are given some weights λ (we assume a trained classifier for now)

Feature	Weights for Class NN	Weights for Class V
PreviousWord=at	0.35	-0.2
CurrentWord=Grace	0.7	0.05
NextWord=Road	0.4	0.3

Maximum Entropy Model: Example

Feature	Weights for Class NN	Weights for Class V
PreviousWord=at	0.35	-0.2
CurrentWord=Grace	0.7	0.05
NextWord=Road	0.4	0.3

- Y= NN and x =“at Grace Road”:

$$p(NN|x) = \frac{\exp(0.35 + 0.7 + 0.4)}{\sum_{y' \in Y} \exp(\sum_{i=1} \lambda_i f_i(x, y'))}$$

Maximum Entropy Model: Example

Feature	Weights for Class NN	Weights for Class V
PreviousWord=at	0.35	-0.2
CurrentWord=Grace	0.7	0.05
NextWord=Road	0.4	0.3

- $Y = V$ and $x = \text{"at } \underline{\text{Grace}} \text{ Road"}$:

$$p(V|x) = \frac{\exp(-0.2 + 0.05 + 0.3)}{\sum_{y' \in Y} \exp(\sum_{i=1} \lambda_i f_i(x, y'))}$$

Maximum Entropy Model: Example

- Denominator:

$$\begin{aligned} & \sum_{y' \in \{NN, V\}} \exp\left(\sum_{i=1} \lambda_i f_i(x, y')\right) \\ &= \exp(0.35 + 0.7 + 0.4) + \exp(-0.2 + 0.05 + 0.3) \\ &= \exp(1.45) + \exp(0.15) \end{aligned}$$

Maximum Entropy Model: Example

$$p(NN|x) = \frac{\exp(0.35 + 0.7 + 0.4)}{\exp(1.45) + \exp(0.15)} = 0.78$$

$$p(V|x) = \frac{\exp(-0.2 + 0.05 + 0.3)}{\exp(1.45) + \exp(0.15)} = 0.22$$

→ We would classify it as a noun

Maximum Entropy Model: Example

- A different view:

$$p(V|x) = \frac{\text{score}(x, V)}{\text{score}(x, V) + \text{score}(x, NN)}$$

- ➔ In essence the weights show the importance of a feature and an instance gets the score as the sum of all active features
- ➔ The denominator is just there to normalize our distribution

Maximum Entropy Models

Parameter Learning

Roadmap

- We have found a family of functions „Exponential Family“, which has some parameters λ_i
- We can now choose an appropriate Loss function
 - Then form the gradient of the loss with respect to every λ_i
 - Use Gradient Descent to optimize the parameters

Maximum Entropy: Parameter Learning

- So far we derived the classifier, but we left how we determine the λ 's
- In fact in order to „train“ the classifier, we need:
 - some data D consisting of tuples (x, y)
 - A **loss function** (which tells us which lambda values are better than others)
- This classifier is usually trained using „**Maximum Likelihood**“

Maximum Entropy: Parameter Learning

- We can do this, using “Maximum Likelihood”:

$$p(y|D, \lambda) = \prod_{(x,y) \in D} p(y|x)$$

- We take the logarithm (numerical stability)

$$p(y|D, \lambda) = \sum_{(x,y) \in D} \log p(y|x)$$

→ Maximum, when the model always predicts the correct class y with $p(y|x)=1$

Maximum Entropy Model: Parameter Learning

- Taking the log, for easier computations (“Log Likelihood”):

$$\begin{aligned}
 p(y|D, \lambda) &= \sum_{(x,y) \in D} \log p(y|x) \\
 &= \sum_{(x,y) \in D} \log \frac{\exp(\sum_{i=1} \lambda_i f_i(x, y))}{\sum_{y' \in Y} \exp(\sum_{i=1} \lambda_i f_i(x, y'))} \\
 &= \sum_{(x,y) \in D} \log \exp \left(\sum_{i=1} \lambda_i f_i(x, y) \right) - \log \sum_{y' \in Y} \exp \left(\sum_{i=1} \lambda_i f_i(x, y') \right)
 \end{aligned}$$

- ➔ find parameter so that this is maximized!
- ➔ Equivalently you can minimize the negative log likelihood

Maximum Entropy Model: Parameter Learning

- In order to be able to optimize the parameters using Gradient Descent, we still need the gradient...
- Objective function:

$$\sum_{(x,y) \in D} \underbrace{\log \exp \left(\sum_{i=1} \lambda_i f_i(x, y) \right)}_A - \log \sum_{y' \in Y} \underbrace{\exp \left(\sum_{i=1} \lambda_i f_i(x, y') \right)}_B$$

The Derivative I: Numerator (A)

$$\begin{aligned}\frac{\partial A}{\partial \lambda_i} &= \frac{\partial \sum_{x,y} \log e^{\sum_i \lambda_i f_i(x,y)}}{\partial \lambda_i} \\ &= \frac{\partial \sum_{x,y} \sum_i \lambda_i f_i(x,y)}{\partial \lambda_i} \\ &= \sum_{x,y} f_i(x,y)\end{aligned}$$

→ Derivative of the numerator is: the empirical count(f_i, y)

The Derivative II: Denominator

$$\begin{aligned}
 \frac{\partial B}{\partial \lambda_i} &= \frac{\partial \sum_{x,y} \log \sum_{y'} e^{\sum_i \lambda_i f_i(y',x)}}{\partial \lambda_i} \\
 &= \sum_{x,y} \frac{1}{\sum_{y''} e^{\sum_i \lambda_i f_i(y'',x)}} \cdot \frac{\partial \sum_{y'} e^{\sum_i \lambda_i f_i(y',x)}}{\partial \lambda_i} \\
 &= \sum_{y,x} \frac{1}{\sum_{y''} e^{\sum_i \lambda_i f_i(y'',x)}} \cdot \sum_{y'} e^{\sum_i \lambda_i f_i(y',x)} \frac{\partial \sum_i \lambda_i f_i(y',x)}{\partial \lambda_i} \\
 &= \sum_{y,x} \sum_{y'} \frac{e^{\sum_i \lambda_i f_i(y',x)}}{\sum_{y''} e^{\sum_i \lambda_i f_i(y'',x)}} \cdot \frac{\partial \sum_i \lambda_i f_i(y',x)}{\partial \lambda_i} \\
 &= \sum_{x,y} \sum_{y'} p(y'|x, \lambda) \cdot f_i(y', x) = \text{predicted count}(f_i, \lambda)
 \end{aligned}$$

The Derivative III

$$\frac{\partial \log P(C | D, \lambda)}{\partial \lambda_i} = \text{actual count}(f_i, C) - \text{predicted count}(f_i, \lambda)$$

- The optimum parameters are the ones for which each feature's **predicted expectation** equals its **empirical expectation**. The optimum distribution is:
 - Always unique (but parameters may not be unique)
 - Always exists (if feature counts are from actual data).

Recap: Maximum Entropy

- We started with the definition of Entropy
 - Went through some ugly math to find our family $p(y|x)$
 - Learned how to optimize the parameters λ_i of our model
 - Learned how to apply the model to new data
- ➔ We could now use a Maximum Entropy classifier to derive our node scores! And we can integrate arbitrary features!

Example: Maximum Entropy

- Running example:
 - Determine the POS-Tags of the sentence

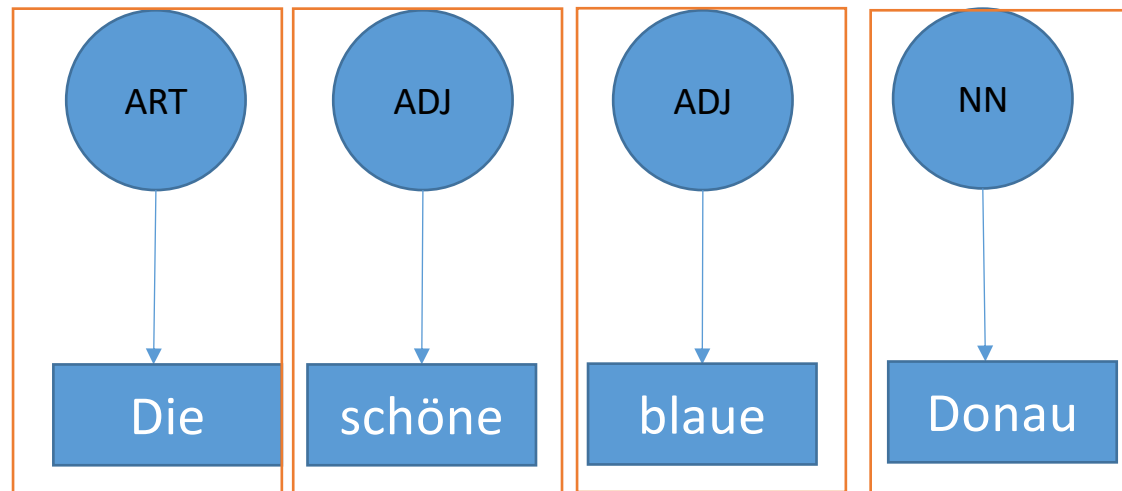
Die schöne blaue Donau

The beautiful blue Danube

- Available POS-Tags: {ART,ADJ,NN}

Example: Maximum Entropy


- We use a local MaxEnt at every token



The beautiful blue Danube

- ➔ 4 local classifier (we assume it is always the same classifier, reused)
- ➔ Only available feature template $f(\mathbf{x}, y_i)$
 - That is, features can be calculated from the entire word sequence and the label at the current position

Example: MaxEnt

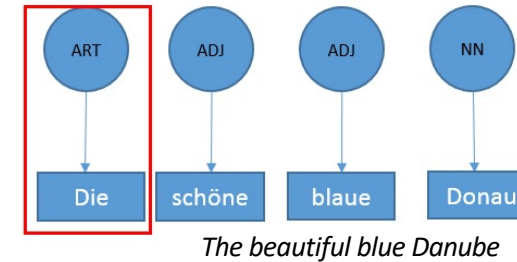


The beautiful blue Danube

- Given example features and their weights, trained on an arbitrary data set

Feature	λ for ART	λ for ADJ	λ for NN
CurrentWord=Die	0.6	0.1	0.25
CurrentWord=schöne	-0.1	0.8	0.3
CurrentWord=blaue	0.1	0.6	0.2
CurrentWord=Donau	0.1	0.1	1.4

Example: MaxEnt



Feature	λ for ART	λ for ADJ	λ for NN
CurrentWord=Die	0.6	0.1	0.25
CurrentWord=schöne	-0.1	0.8	0.3
CurrentWord=blaue	0.1	0.6	0.2
CurrentWord=Donau	0.1	0.1	1.4


$$p(ART|x) = \frac{\exp(0.6)}{\exp(0.6) + \exp(0.1) + \exp(0.25)} = 0.44$$

$$p(ADJ|x) = \frac{\exp(0.1)}{\exp(0.6) + \exp(0.1) + \exp(0.25)} = 0.26$$

$$p(NN|x) = \frac{\exp(0.25)}{\exp(0.6) + \exp(0.1) + \exp(0.25)} = 0.30$$

→ "Die" is an ART

Example: MaxEnt



The beautiful blue Danube

Feature	λ for ART	λ for ADJ	λ for NN
CurrentWord=Die	0.6	0.1	0.25
CurrentWord=schöne	-0.1	0.8	0.3
CurrentWord=blaue	0.1	0.6	0.2
CurrentWord=Donau	0.1	0.1	1.4


$$p(ART|x) = \frac{\exp(-0.1)}{\exp(-0.1) + \exp(0.8) + \exp(0.3)} = 0.2$$

$$p(ADJ|x) = \frac{\exp(0.8)}{\exp(-0.1) + \exp(0.8) + \exp(0.3)} = 0.5$$

$$p(NN|x) = \frac{\exp(0.3)}{\exp(-0.1) + \exp(0.8) + \exp(0.3)} = 0.3$$

➔ "schöne" is an ADJ

Example: MaxEnt



Feature	λ for ART	λ for ADJ	λ for NN
CurrentWord=Die	0.6	0.1	0.25
CurrentWord=schöne	-0.1	0.8	0.3
CurrentWord=blaue	0.1	0.6	0.2
CurrentWord=Donau	0.1	0.1	1.4


$$p(ART|x) = \frac{\exp(0.1)}{\exp(0.1) + \exp(0.6) + \exp(0.2)} = 0.26$$

$$p(ADJ|x) = \frac{\exp(0.6)}{\exp(0.1) + \exp(0.6) + \exp(0.2)} = 0.44$$

$$p(NN|x) = \frac{\exp(0.2)}{\exp(0.1) + \exp(0.6) + \exp(0.2)} = 0.3$$

→ "blaue" is an ADJ

Example: MaxEnt



The beautiful blue Danube

Feature	λ for ART	λ for ADJ	λ for NN
CurrentWord=Die	0.6	0.1	0.25
CurrentWord=schöne	-0.1	0.8	0.3
CurrentWord=blaue	0.1	0.6	0.2
CurrentWord=Donau	0.1	0.1	1.4

$$p(ART|x) = \frac{\exp(0.1)}{\exp(0.1) + \exp(0.1) + \exp(1.4)} = 0.176$$

$$p(ADJ|x) = \frac{\exp(0.1)}{\exp(0.1) + \exp(0.1) + \exp(1.4)} = 0.176$$

$$p(NN|x) = \frac{\exp(1.4)}{\exp(0.1) + \exp(0.1) + \exp(1.4)} = 0.647$$

➔ "Donau" is a NN

Maximum Entropy: Recap

- Only local information is used
 - ➔ Decoding is always greedy
- Only one feature template $f(\mathbf{x}, y_i)$
- Still manages to classify the sequence correct, because the features are expressive enough
- In general the used features do not generalize to unseen words