

# Coreference Resolution

Machine Learning Approaches

# What is Coreference Resolution?

## Example

[Barack Obama]<sub>1</sub><sup>1</sup> nominated [Hillary Clinton]<sub>2</sub><sup>2</sup> as [[his]<sub>3</sub><sup>1</sup> secretary of state]<sub>4</sub><sup>3</sup> on [Monday]<sub>5</sub><sup>4</sup>. [He]<sub>6</sub><sup>1</sup>

- Superscript: ID of an entity
- Subscript: ID of a mention

# Coreference as Clustering

- Bell number:

Number of Mentions	Bell number
10	115975
20	51724158236496
30	846749014529889671069667
50	$1.8572414972124E+47$
100	$2.3 \cdot 10^{117}$

# Coreference as Clustering

- Conditional Random Fields for Coreference:

$$p(y|x) = \frac{\exp(\sum_{features} \lambda_f f(x, y))}{\sum_{\hat{y} \in clusterings} \exp(\sum_{features} \lambda_f f(x, \hat{y}))}$$

Diagram illustrating the components of the equation:

- Arbitrary Clustering**: Points to the variable  $y$  in the numerator.
- Input text**: Points to the variable  $x$  in both the numerator and denominator.
- Features for the Clustering  $y$** : Points to the  $f(x, y)$  term in the numerator.

➔ Infeasible to calculate, so we need to go back to (crude) approximations!


# Coreference as Clustering

- Global methods are infeasible, since there is no dynamic programming for the scoring of the clusters
- ➔ We are looking for approaches, that search for local solutions, which we can combine into a clustering



# Coreference as Dependency Parsing

- Let us compare Coreference Resolution to Dependency Parsing



- Instead of tokens, which are looking for a head, we are searching a head for every mention
- Newly introduced mentions will be resolved to the artificial *root node*

# Coreference as Dependency Parsing

- The problem has now become much simpler:
  - For every mention, we have to find any antecedent
  - First appearances of an entity will be mapped to root
- ➔ Instead of predicting an entire cluster at once, we will now only score **links between mentions**
- This results in at least the following approaches:
  1. Mention-Pair Modell
  2. Mention-Ranking Modell
  3. Entity-Mention Modell
  4. Cluster-Ranking Modell

# Machine Learning: Mention-Pair

- Given a labelled corpus:

«Das ist's, <sup>♂</sup>Richard<sup>2</sup> – ja, unklar geblieben ist <sup>♀</sup>mir<sup>325</sup> Manches, schon früher, aber <sup>♀</sup>ich<sup>325</sup> habe es immer  
<sup>♀</sup>meiner<sup>325</sup> eigenen Unwissenheit zugeschrieben; und da es meistens Dinge betraf, die <sup>♀</sup>ich<sup>325</sup> nicht gut – die <sup>♀</sup>ich<sup>325</sup>  
möglicherweise ganz falsch verstanden – um die <sup>♀</sup>ich<sup>3</sup> <sup>♂</sup>Dich<sup>2</sup> nicht fragen wollte, aus –«<sup>325</sup>

Ausschnitt aus: Aston Louise: Lydia

- Let us now model this into a learnable problem



# Machine Learning: Mention-Pair

- We first extract all instances  $(x, y)$  as follows:
  - Every  $x$  is now a pair of mentions  $(m_i, m_j)$
  - Every  $y$  is a label of  $\{\text{coref}, \text{non-coref}\}$

➔ Simple binary, unstructured classification

- And on top, we get a quadratic amount of instances
  - E.g. for the labelled „Aston Louise- Lydia“ corpus:  $\frac{643 \cdot 642}{2} = 206403$  !!

# Machine Learning: Mention-Pair

- Is getting so many training examples a good thing?
  - Laws et al: of 1,7 million instances, 98.5% were non-coref!  
(Data Skewness)
- ➔ Either very robust classifier are required or a strategy to „sample“ the instances which carry most information:
  1. Sampling
  2. Active-Learning

# Data Skewness

- Why is **Data Skewness** even a problem?
- Example with Maximum-Entropy:
  - Usually a good feature will receive a high weight  $\lambda$  during training
  - Being a good feature means, that is appears more often with class A, than with class B
  - But lets assume the calculation of our features is noisy (e.g. through different preprocessing or a noisy data set)
    - ➔ If the feature is misleading just a single time out of 98 instances, then its entire expressiveness is already gone, because of the imbalance in our data set!
- ➔ Getting a more or less equal amount of instances for each class is a desired property!

# Data Skewness

- Why is **Data Skewness** even a problem?
  - Example with Perceptron:
    - You remember the bias  **$b$**
    - If there is such a large skewness in your data, you can think of your classification process as follows:  
Each instance has to persuade the classifier not to be of the most frequent class!  
(otherwise all the score will be negated by adding the bias!)
- ➔ Getting a more or less equal amount of instances for each class is a desired property!

# Machine Learning: Mention-Pair

- Method of Soon et al. (2001):
  - Instead of building all pairs  $(m_i, m_j)$  of the training data, just take:

**For each** mention  $m_i$ :

Iterate backwards in the text,  $(m_j, j < i)$

Insert the pair  $(m_i, m_j)$  into the train data

**If**  $(m_i, m_j)$  co-refer,

**Break;**

- In words:  
For every mention  $m_i$  only add pairs with other mentions until you find the first mention that co-refers.

# Machine Learning: Mention-Pair

- Method of Soon et al. (2001):

## Example

[Barack Obama]<sub>1</sub> nominated [Hillary Clinton]<sub>2</sub> as [[his]<sub>3</sub> secretary of state]<sub>4</sub> on [Monday]<sub>5</sub>. [He]<sub>6</sub>

- This method produces the following pairs:
  - (Hillary Clinton, Barack Obama): Non-Coref
  - (his, Hillary Clinton): Non-Coref
  - (his, Barack Obama): Coref
  - (secretary of state, *X*): Non-Coref (3 Instances)
  - (monday, *X*): Non-Coref (4 Instances)
  - (He, Monday): Non-Coref
  - (He, secretary of state): Non-Coref
  - (He, his): Coref

13 instances  
(11 vs. 2)

# Machine Learning: Mention-Pair

- Features for our classifier:
  - E.g. for the pair (**Barack Obama**, **his**)
  - [**Barack Obama**]<sub>1</sub><sup>1</sup> nominated [Hillary Clinton]<sub>2</sub><sup>2</sup> as [[**his**]<sub>3</sub><sup>1</sup> secretary of state]<sub>4</sub><sup>3</sup> on [Monday]<sub>5</sub><sup>4</sup>. [He]<sub>6</sub><sup>1</sup>
- Brainstorming: ...  
(you could think about this yourself, this is the key issue as of why this task is so hard!)

# Machine Learning: Mention-Pair

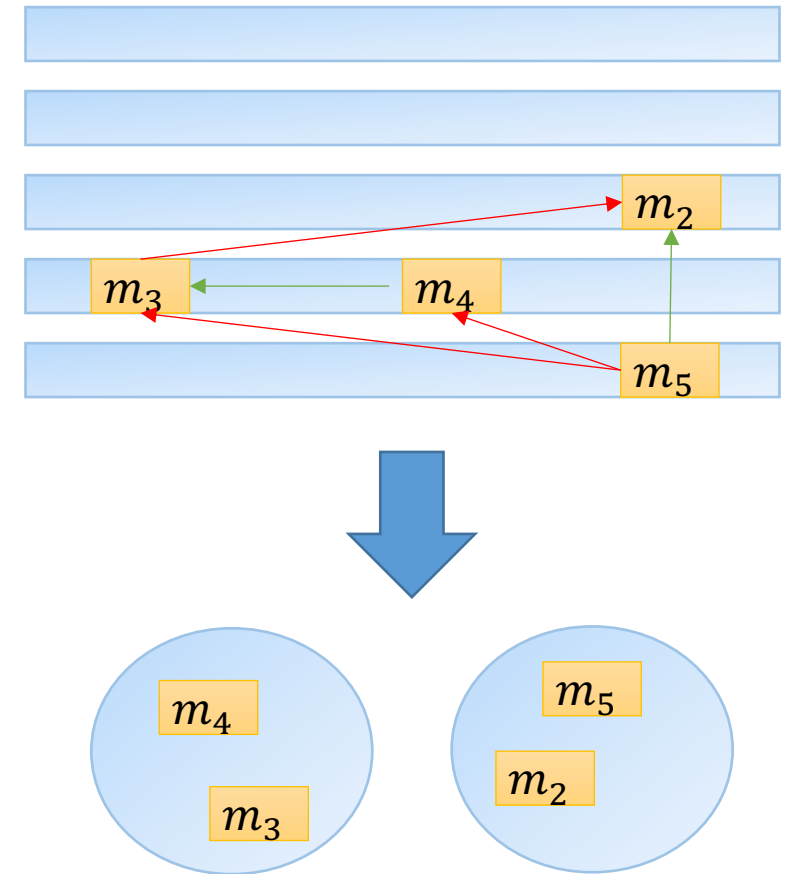
- Features of Soon et al.

Type	Description	Value
Distance	Distance between $m_i$ and $m_j$ in terms of the number of sentences	Integer
String-matching	Is $m_j$ an alias of $m_i$ ?	Boolean
	Do $m_i$ and $m_j$ match after stripping of articles and demonstrative pronouns?	Boolean
Grammatical	Does $m_j$ start with an definite article?	Boolean
	Does $m_j$ start with a demonstrative pronoun?	Boolean
	Is $m_i$ pronominal?	Boolean
	Is $m_j$ pronominal?	Boolean
	Do $m_i$ and $m_j$ agree in number?	Boolean
	Do $m_i$ and $m_j$ agree in gender?	Boolean
	Do $m_i$ and $m_j$ both contain a proper name?	Boolean
Syntactic	Is $m_j$ an apposition?	Boolean
Semantic	Do $m_i$ and $m_j$ agree in semantic class	Boolean



# Machine Learning: Mention-Pair

- Application of the approach (according to Soon):
  - Create instances „on-the-fly“ and predict them
    - Only backwards in the text
  - Stop as soon as you get the first **Coref** prediction.  
Add the edge to the solution
- ➔ Build the clustering from the positive „edges“




# Machine Learning: Mention-Pair


- Problems:

- Mention-Pair produces inconsistencies

- $(m_3, m_2) \Rightarrow$  not coref
- $(m_2, m_1) \Rightarrow$  coref
- $(m_3, m_1) \Rightarrow$  coref



- Mentions are not competing (instances are independent)



- Local features: They can only take two mentions into account!


# Machine Learning: Mention-Ranking

- For every mention, „rank“ all candidates against each other and finally take the best one



# Machine Learning: Mention-Ranking

- Ranking instead of classification: But how?
- Idea: No complete ranking is required, all we need is a correct first position!
- If for example we are looking for the best candidate for  $m_5$ 
  - We have 4 mentions competing:  $m_4$  up to  $m_1$




# Mention-Ranking: Tournament Model

- The tournament model is a determines a ranking with a simple binary classifier
  - And all we need to adjust is the way we model our instances
  - So far our instances have been of the form:  $(x = (m_i, m_j), y \in \{\text{Coref}, \text{Non-Coref}\})$
  - The tournament model changes this into:  $(x = (m_i, m_j), (m_i, m_k), y \in \{\text{First-pair}, \text{Second-pair}\})$
  - So the classifier gets two pairs and decides which pair is the better one
  - Classifier is still binary in nature
    - First-pair → We prefer the first pairing  $((m_5, m_4); (m_5, m_3))$
    - Second-pair → We prefer the second pairing  $((m_5, m_4); (m_5, m_3))$
- ➔ It is just a smart way to apply a binary classifier!

# Mention-Ranking: Tournament Model

- How to apply it:
  - We always have two pairings with one „defeating“ the other
  - Once we have reached the first mention, there is a clear winner
  - E.g. (we want to find the the antecedent for  $m_5$ ):



$(m_5, m_4)$  vs.  $(m_5, m_3) \rightarrow (m_5, m_4)$

$(m_5, m_4)$  vs.  $(m_5, m_2) \rightarrow (m_5, m_2)$

$(m_5, m_2)$  vs.  $(m_5, m_1) \rightarrow (m_5, m_2)$

$\rightarrow m_5$  is resolved to  $m_2$

# Mention-Ranking: Tournament Model

- Remaining problems:

1. Mention-Ranking produces inconsistencies

- $(m_3, m_2) \Rightarrow \text{Not-Coref}$
- $(m_2, m_1) \Rightarrow \text{Coref}$
- $(m_3, m_1) \Rightarrow \text{Coref}$

2. ~~Mentions compete against each other~~



➔ This is solved at least partially!

3. Local Features

- Features can yet again only be calculated on the basis of a single edge


4. How to recognize the first appearances? („discourse-new“)

- There is always one pairing that remains ...



# Machine Learning: Entity-Mention


- Let us improve on the problem of local features, by:
  - Creating the clusters in an iterative fashion
  - E.g. Luo 2006 using the Bell-Tree:





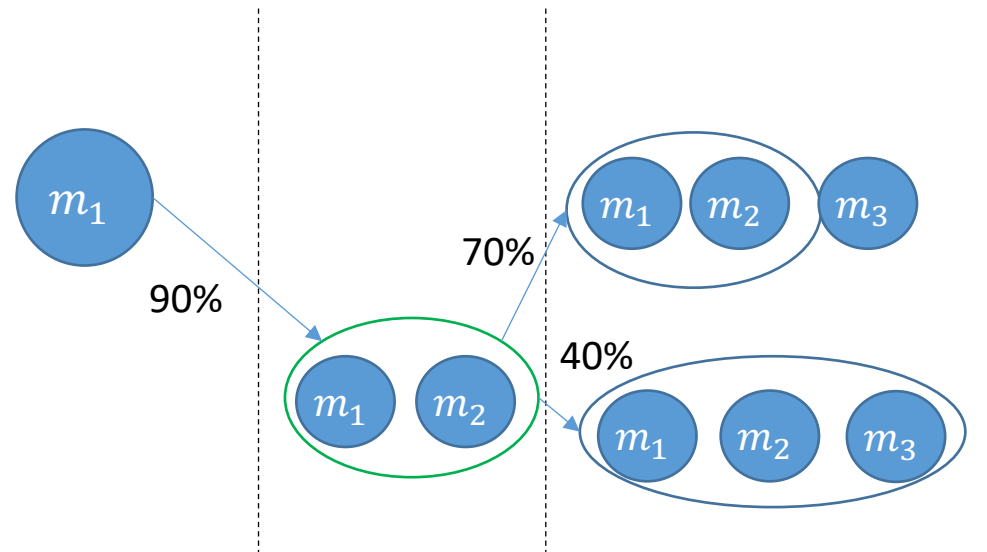
# Machine Learning: Entity-Mention

- Let us improve on the problem of local features, by:
  - Creating the clusters in an iterative fashion
  - E.g. Luo 2006 using the Bell-Tree:




# Machine Learning: Entity-Mention

- Let us improve on the problem of local features, by:
  - Creating the clusters in an iterative fashion
  - E.g. Luo 2006 using the Bell-Tree:



# Machine Learning: Entity-Mention

- Let us improve on the problem of local features, by:
  - Creating the clusters in an iterative fashion
  - E.g. Luo 2006 using the Bell-Tree:



# Machine Learning: Entity-Mention

- Clusters are now built in an iterative fashion, mentions are resolved from the beginning of a text up to the end of a text (similar as to how a human would do it)
- We will always only keep the best solution
  - ➔ State-space shrinks to a manageable size!
- Instances will be  $(cluster_i, m_k)$ 
  - So we can now access features of the (partial) cluster and a mention!

# Machine Learning: Entity-Mention

- Remaining problems:

- ~~1. Mention-Ranking produces inconsistencies~~


- $(m_3, m_2) \Rightarrow \text{Not-Coref}$
- $(m_2, m_1) \Rightarrow \text{Coref}$
- $(m_3, m_1) \Rightarrow \text{Coref}$

2. Mentions do not compete against each other!

- ~~3. Local Features~~


4. How to recognize the first appearances? („Discourse-new“)

- There is always one pairing that remains ...




# Machine Learning: Cluster-Ranking

- Combine the strengths of the **Mention ranking** model with the strengths of the **Entity-mention** model
- For example by using a hierarchical clustering (e.g. HAC)
- Start with all mentions being their own „cluster“



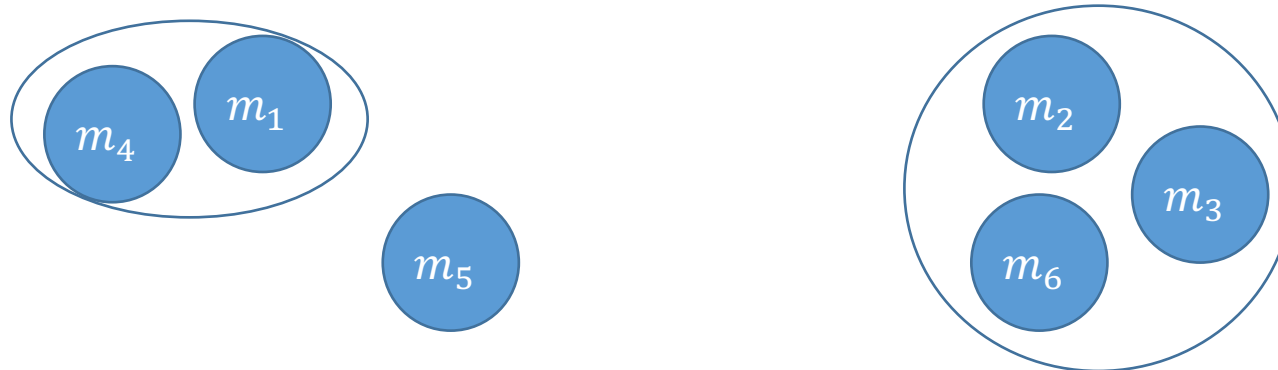
# Machine Learning: Cluster-Ranking

- Start with all mentions being their own „cluster“
- Use the tournament model to get the best pairing for each mention
  - ➔ But only apply the one which is the best according to some criteria



# Machine Learning: Cluster-Ranking

- Start with all mentions being their own „cluster“
- Use the tournament model to get the best pairing for each mention
  - But only apply the one which is the best according to some criteria
- Repeat until there is no more combination that exceeds a certain threshold





# Machine Learning: Cluster-Ranking

- Remaining Problems:
  - ~~1. We produce inconsistencies~~
  - ~~2. Mentions compete against each other~~
  - ~~3. Local Features~~
    - ➔ We can now calculate features using pairs of clusters!
  4. How to recognize the first appearances? („Discourse-new“)
    - There is always one pairing that remains ...

There is one issue  
remaining!

# Machine Learning: Anaphoricity

- The decision, whether a mention has any antecedent in the text is called „Anaphoricity“-Problem
- All ranking approaches need this kind of information!
- Two typical ways to solve this:
  1. Create a dummy mention: if the dummy wins the tournament, then we do not assign any antecedent! (**Joint-Approach J**)
  2. A separate classifier decides, whether a mention is resolved at all! (**Pipeline-Approach P**)



# Recap: Machine Learning for Coreference

- In this lecture we presented different approaches to keep the problem of Coreference Resolution tractable
  - We did this by finding an analogy to the task of Dependency Parsing
- This resulted in four models, that are independent of the classifier that is used:
  - Mention-Pair
  - Mention-Ranking
  - Entity-Pair
  - Cluster-Ranking

Even though the more complicated models tend to outperform the more basic models, there is no solution to Coreference Resolution as of now!

# Evaluation

- Coreference Resolution produces a „clustering“
- For evaluating clusterings, measures such as the MUC metric are addressed in detail in the **chapter "Evaluation"**